

Wenting Zhao

Cornell Tech, New York

wzhao@cs.cornell.edu | [wenting-zhao.github.io](https://github.com/wenting-zhao) | [wenting-zhao](https://www.linkedin.com/in/wenting-zhao)

Education

Cornell University

PH.D., COMPUTER SCIENCE

• Advisors: Claire Cardie, Alexander Rush

New York, New York

Aug 2018 - May 2025 (expected)

Cornell University

M.Sc., COMPUTER SCIENCE

Ithaca, New York

Dec 2022

Worcester Colledge, University of Oxford

VISITING STUDENT

• Advisor: Peter Jeavons

Oxford, UK

Sep 2016 - Jun 2017

Illinois Wesleyan University

B.Sc., COMPUTER SCIENCE

• Honors thesis: Encoding Lexicographical Ordering Constraints in SAT

Bloomington, Illinois

Aug 2014 - Dec 2017

Employment

Jan 2023 - Jun 2024 **Research Intern**, Allen Institute for AI
Aug 2021 - Jan 2022 **Applied Scientist Intern**, Amazon Alexa AI
May 2021 - Aug 2021 **Research Intern**, IBM AI Research

Awards

2024	Rising Star in Generative AI , University of Massachusetts Amherst	
2023	Outstanding Intern Award , Allen Institute for AI	\$ 10,000
2018	Grace Hopper Celebration Travel Grant , Cornell University	\$ 1,500
2015	Eckley Research Fellowship , Illinois Wesleyan University	\$ 4,000
2014	Alumni Scholarship , Illinois Wesleyan University	\$ 64,000

Media Coverage

Aug 14, 2024 [Study suggests that even the best AI models hallucinate a bunch](#), TechCrunch
Aug 4, 2024 [What do people really ask chatbots? It's a lot of sex and homework](#), The Washington Post

Publications

PEER-REVIEWED CONFERENCE/JOURNAL PAPERS

- [C15] **Wenting Zhao**, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, Yuntian Deng. "WildChat: 1M ChatGPT Interaction Logs in the Wild." In: *Proceedings of ICLR*. 2024. (**Spotlight**) [[pdf](#)]
- [C14] John Morris, **Wenting Zhao**, Justin Chiu, Vitaly Shmatikov, Alexander M Rush. "Language Model Inversion." In: *Proceedings of ICLR*. 2024. [[pdf](#)]
- [C13] **Wenting Zhao**, Justin Chiu, Jena Hwang, Faeze Brahman, Jack Hessel, Sanjiban Choudhury, Yejin Choi, Lorraine Xiang Li, Alane Suhr. "UNcommonsense Reasoning: Abductive Reasoning about Uncommon Situations." In: *Proceedings of NAACL*. 2024. [[pdf](#)]

- [C12] **Wenting Zhao**, Ge Gao, Claire Cardie, Alexander M Rush. “*I Could’ve Asked That*: Reformulating Unanswerable Questions.” In: *Proceedings of EMNLP*. 2024. [[pdf](#)]
- [C11] Yuntian Deng, **Wenting Zhao**, Jack Hessel, Xiang Ren, Claire Cardie, Yejin Choi. “WildVis: Open Source Visualizer for Million-Scale Chat Logs in the Wild.” In: *Proceedings of EMNLP (demo)*. 2024. [[pdf](#)]
- [C10] Huihan Li, Yuting Ning, Zeyi Liao, Siyuan Wang, Xiang Lorraine Li, Ximing Lu, **Wenting Zhao**, Faeze Brahman, Yejin Choi, Xiang Ren. “In Search of the Long-Tail: Systematic Generation of Long-Tail Knowledge via Logical Rule Guided Search.” In: *Proceedings of EMNLP*. 2024. [[pdf](#)]
- [C9] **Wenting Zhao**, Justin Chiu, Alexander M Rush, Claire Cardie. “Hop, Union, Generate: Explainable Multi-hop Reasoning without Rationale Supervision.” In: *Proceedings of EMNLP*. 2023. [[pdf](#)]
- [C8] Justin Chiu, **Wenting Zhao**, Derek Chen, Saujas Vaduguru, Alexander Rush, Daniel Fried. “Symbolic Planning and Code Generation for Grounded Dialogue.” In: *Proceedings of EMNLP*. 2023. [[pdf](#)]
- [C7] **Wenting Zhao**, Justin Chiu, Alexander M Rush, Claire Cardie. “Abductive Commonsense Reasoning Exploiting Mutually Exclusive Explanations.” In: *Proceedings of ACL*. 2023. [[pdf](#)]
- [C6] **Wenting Zhao**, Weiqi Sun, Konstantine Arkoudas, Claire Cardie. “Compositional Task-Oriented Parsing as Abstractive Question Answering.” In: *Proceedings of NAACL*. 2022. [[pdf](#)]
- [C5] Di Chen, Yiwei Bai, Sebastian Ament, **Wenting Zhao**, Dan Guevarra, Lan Zhou, Bart Selman, R. Bruce van Dover, John M. Gregoire, Carla P. Gomes. Automating crystal-structure phase mapping by combining deep learning with constraint reasoning. In: *Nature Machine Intelligence*. 2021. (**Cover story**) [[link](#)]
- [C4] **Wenting Zhao**, Shufeng Kong, Junwen Bai, Daniel Fink, and Carla Gome. “HOT-VAE: Learning High-Order Label Correlation for Multi-Label Classification via Attention-Based Variational Autoencoders” In: *Proceedings of AAAI (Computational Sustainability)*. 2021. [[pdf](#)]
- [C3] Di Chen, Yiwei Bai, **Wenting Zhao**, Sebastian Ament, John Gregoire, and Carla Gomes. “Deep Reasoning Networks for Unsupervised Pattern De-mixing with Constraint Reasoning.” In: *Proceedings of ICML*. 2020. [[pdf](#)]
- [C2] **Wenting Zhao**, Mark H. Liffiton, Peter G. Jeavons, Dan Roberts. “Finding graph decompositions via SAT”. In: *Proceedings of ICTAI*. 2017. [[pdf](#)]
- [C1] **Wenting Zhao**, Mark H. Liffiton. Parallelizing Partial MUS Enumeration. In: *Proceedings of ICTAI*. 2016. [[pdf](#)]

TUTORIALS

- [T1] **Wenting Zhao**, Mor Geva, Bill Yuchen Lin, Michihiro Yasunaga, Aman Madaan, Tao Yu. “Complex Reasoning in Natural Language.” In: *Proceedings of ACL*. 2023. [[link](#)]

PEER-REVIEWED WORKSHOP PAPERS

- [W2] **Wenting Zhao**, Alexander M Rush, Claire Cardie. “Commonsense Reasoning for Question Answering with Explanations.” In: *Workshop on Commonsense Representation and Reasoning @ ACL (Non-archival)*. 2022 [[link](#)]
- [W1] Justin T Chiu, **Wenting Zhao**, Daniel Fried, Alexander M Rush. “Modeling Perspective-Dependent Ambiguity in Collaborative Dialogue.” In: *The Third Workshop on Wordplay: When Language Meets Games @ NAACL*. 2022. [[link](#)]

PREPRINTS

- [P8] **Wenting Zhao**, Nan Jiang, Celine Lee, Justin T Chiu, Claire Cardie, Matthias Gallé, Alexander M Rush. “Commit0: Library Generation from Scratch.” 2024. [[pdf](#)]
- [P7] **Wenting Zhao**, Alexander M Rush, Tanya Goyal. 2024. “Challenges in Trustworthy Human Evaluation of Chatbots”
- [P6] **Wenting Zhao**, Tanya Goyal, Yu Ying Chiu, Liwei Jiang, Benjamin Newman, Abhilasha Ravichander, Khyathi Chandu, Ronan Le Bras, Claire Cardie, Yuntian Deng, Yejin Choi. “WildHallucinations: Evaluating Long-form Factuality in LLMs with Real-World Entity Queries.” 2024. [[pdf](#)]
- [P5] Shangyi Geng, **Wenting Zhao**, Alexander M Rush. “Great Memory, Shallow Reasoning: Limits of k NN-LMs.” 2024. [[pdf](#)]

- [P4] Yi Lu, Jing Nathan Yan, Songlin Yang, Justin T Chiu, Siyu Ren, Fei Yuan, **Wenting Zhao**, Zhiyong Wu, Alexander M Rush. “A Controlled Study on Long Context Extension and Generalization in LLMs.” 2024. [\[pdf\]](#)
- [P3] Shachar Don-Yehiya, Ben Burtenshaw, Ramon Fernandez Astudillo, Cailean Osborne, Mimansa Jaiswal, Tzu-Sheng Kuo, **Wenting Zhao**, Idan Shenfeld, Andi Peng, Mikhail Yurochkin, Atoosa Kasirzadeh, Yangsibo Huang, Tatsunori Hashimoto, Yacine Jernite, Daniel Vila-Suero, Omri Abend, Jennifer Ding, Sara Hooker, Hannah Rose Kirk, Leshem Choshen. “The Future of Open Human Feedback.” 2024. [\[pdf\]](#)
- [P2] Shaden Shaar*, Wayne Chen*, Maitreyi Chatterjee, Barry Wang, **Wenting Zhao**, Claire Cardie. “Are *Triggers* Needed for Document-Level Event Extraction?” 2024.
- [P1] **Wenting Zhao**, Ibrahim Abdelaziz, Julian Dolby, Kavitha Srinivas, Mossad Helali, Essam Mansour. “Serenity: Library Based Python Code Analysis for Code Completion and Automated Machine Learning.” 2023. [\[pdf\]](#)

Talks

Dec 2024 **Upcoming talk**, Princeton PLI

Nov 2024 **Upcoming talk**, EPFL

Nov 2024 **Upcoming talk**, Data Science Seminar, University of Queensland

Nov 2024 **Commit0: Library Generation from Scratch**, IBM Research

Sep 2024 **Evaluation and Reasoning in Real-world Scenarios**, Rising Stars in GenAI, UMass Amherst

Feb 2024 **Tutorial: Complex Reasoning over Natural Language**, Cornell NLP Seminar

Nov 2023 **570K ChatGPT Interaction Logs In The Wild**, USC ISI/NLG Seminar

Nov 2023 **570K ChatGPT Interaction Logs In The Wild**, NLP with Friends

Nov 2023 **570K ChatGPT Interaction Logs In The Wild**, NLP Group at UCSB

Oct 2023 **570K ChatGPT Interaction Logs In The Wild**, NLP Group at UCI

Jul 2023 **Abductive Commonsense Reasoning Exploiting Mutually Exclusive Explanations**, ACL

Jul 2022 **Compositional Task-Oriented Parsing as Abstractive Question Answering**, NAACL

Teaching Experience

TUTORIAL CO-INSTRUCTOR

ACL 2023 Complex Reasoning in Natural Language
with Mor Geva, Bill Yuchen Lin, Michihiro Yasunaga, Aman Madaan, Tao Yu

TEACHING ASSISTANT

2020-2021 CS 4700: Foundations of Artificial Intelligence, Cornell University
Instructor: Haym B. Hirsh

2015-2017 CS 127: Computer Science I, Illinois Wesleyan University
Instructors: Mark Liffiton, Kamil Samara

Mentoring

2023-2024 **Shangyi Geng**, MS student, Cornell Tech

2024 **Nan Jiang**, BS student, UCSD

2024 **Deniz Boloni-Turgut**, BS student, Cornell University

2024 **Dhriti Verma**, BS student, Cornell University

2024 **Max Whitton**, BS student, Cornell University

2024 **Jingtian (Peter) Wu**, BS student, Cornell University

2024 **Eugene Choi**, BS student, Cornell University

2024 **Aishanur Aydin**, BS student, Cornell University

2024 **Leo Lu**, BS student, Cornell University

2024 **Alicia Fulbright**, BS student, Northeastern University

Professional Services

WORKSHOP CO-ORGANIZER

ACL 2024 Natural Language Reasoning and Structured Explanations Workshop

REVIEWER/PROGRAM COMMITTEE

Conferences:

- **ACL** (2023)
- **EMNLP** (2022, 2023)
- **ACL Rolling Review** (2022, 2023, 2024)
- **NeurIPS** (2021, 2022, 2023)
- **ICLR** (2021, 2022, 2023)
- **ICML** (2021, 2022, 2023)

COMMUNITY SERVICE

COLM 2024 Student Social Event, Organizer
EMNLP 2023 Reasoning Birds-of-a-Feather, Session Co-chair

DEPARTMENTAL SERVICE

2020 Cornell CS: Teaching Committee Member (for department-wide course planning during Covid-19)
2019 Cornell CS: Co-organizer for CS PhD Visit Day
2018-2021 Cornell CS: Admission Committee Member for the Ph.D. Program